

1. prefix sum+binary search.  $O(n \log n)$ .
2. offline, generate all samples in a batch. reduce to sorting.
3. use the alias method [https://en.wikipedia.org/wiki/Alias\\_method](https://en.wikipedia.org/wiki/Alias_method).  $O(n)$  preprocessing,  $O(1)$  per query.
4. after normalizing the weights, divide the cdf into  $n$  buckets each with length  $\frac{1}{n}$ . when we perform a query, first randomly pick a bucket, then randomly select an element in the bucket by brute force, in time proportional to the number of elements in that bucket. in expectation there are  $O(1)$  elements in the picked bucket, so  $O(1)$  per query in expectation.

remark. some related papers for more complicated query ranges: [2], [1].

## References

- [1] Peyman Afshani and Jeff M Phillips. Independent range sampling, revisited again. *arXiv preprint arXiv:1903.08014*, 2019.
- [2] Peyman Afshani and Zhewei Wei. Independent range sampling, revisited. In *25th Annual European Symposium on Algorithms (ESA 2017)*. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik, 2017.